

Empathize data wrangling and EDA Vis

August 5, 2023

1 Empathize Stage

1.1 Data Wrangling and EDA Visualization

1.1.1 Aavail Project

Necessary libraries:

```
[ ]: import pandas as pd
import numpy as np
import scipy.stats as stats
import matplotlib.pyplot as plt
%matplotlib inline

import seaborn as sns
from pyampute.exploration.md_patterns import mdPatterns
from pyampute.exploration.mcar_statistical_tests import MCARTest
from sklearn.impute import IterativeImputer

sns.set_theme()
```

```
[ ]: df = pd.read_csv('aavail-data-visualization.csv')
df.head()
```

	customer_id	country_name	age	customer_name	is_subscriber	\
0	1	united_states	21.0	Kasen Todd	True	
1	2	singapore	31.0	Ensley Garza	False	
2	3	united_states	22.0	Lillian Carey	False	
3	4	united_states	21.0	Beau Christensen	True	
4	5	singapore	22.0	Ernesto Gibson	True	

	subscriber_type	num_streams
0	aavail_premium	23.0
1	NaN	12.0
2	aavail_premium	22.0
3	aavail_basic	19.0
4	aavail_premium	23.0

Presenting the shape and variables details of the data frame.

```
[ ]: print(f'The shape of the DataFrame: {df.shape}')
      print('-----')
      print(df.info())
```

The shape of the DataFrame: (1000, 7)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1000 entries, 0 to 999

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	customer_id	1000 non-null	int64
1	country_name	1000 non-null	object
2	age	1000 non-null	float64
3	customer_name	1000 non-null	object
4	is_subscriber	1000 non-null	bool
5	subscriber_type	928 non-null	object
6	num_streams	954 non-null	float64

dtypes: bool(1), float64(2), int64(1), object(3)

memory usage: 48.0+ KB

None

We found that there are two variables **subscriber_type**, **num_streams** with *missing data*.

Check the percentage of missing data in each variables.

```
[ ]: missColumns = ['subscriber_type', 'num_streams']
      df[missColumns].isna().mean()
```

subscriber_type 0.072

num_streams 0.046

dtype: float64

We have less than 10% of missing value in each variable.

Check the patterns of Missing Data.

```
[ ]: df[missColumns].isna().value_counts()/1000
```

subscriber_type	num_streams	
False	False	0.883
True	False	0.071
False	True	0.045
True	True	0.001

dtype: float64


We found that there is a pattern so we will plot the patterns of Missing Data as a Map, Trying to figure it out.

```
[ ]: mdp = mdPatterns()
      patterns = mdp.get_patterns(df)
```

```
print(patterns)
```

c:\Users\ASUS\anaconda3\lib\site-packages\pyampute\exploration\md_patterns.py:120: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
group_values = group_values.append(colsums, ignore_index=True)
```



	row_count	customer_id	country_name	age	\
rows_no_missing	883	1	1	1	
1	71	1	1	1	
2	45	1	1	1	
3	1	1	1	1	
n_missing_values_per_col		0	0	0	

	customer_name	is_subscriber	num_streams	\
rows_no_missing	1	1	1	
1	1	1	1	
2	1	1	0	
3	1	1	0	
n_missing_values_per_col	0	0	46	

	subscriber_type	n_missing_values
rows_no_missing	1	0
1	0	1
2	1	1
3	0	2
n_missing_values_per_col	72	118

We found that the missing data have pattern that makes us think that the missing data are not *MCAR*

Some Wrangling and Investigations.

subscriber_type vs country_name

```
[ ]: df.groupby(['subscriber_type'], dropna=False)['country_name'].value_counts().
      ↪to_frame(name = 'Count').tail(2)
```

		Count
subscriber_type	country_name	
NaN	united_states	51
	singapore	21

This makes us think that there is a bias due to **Missing Data**, so let's find the percentage for each one.

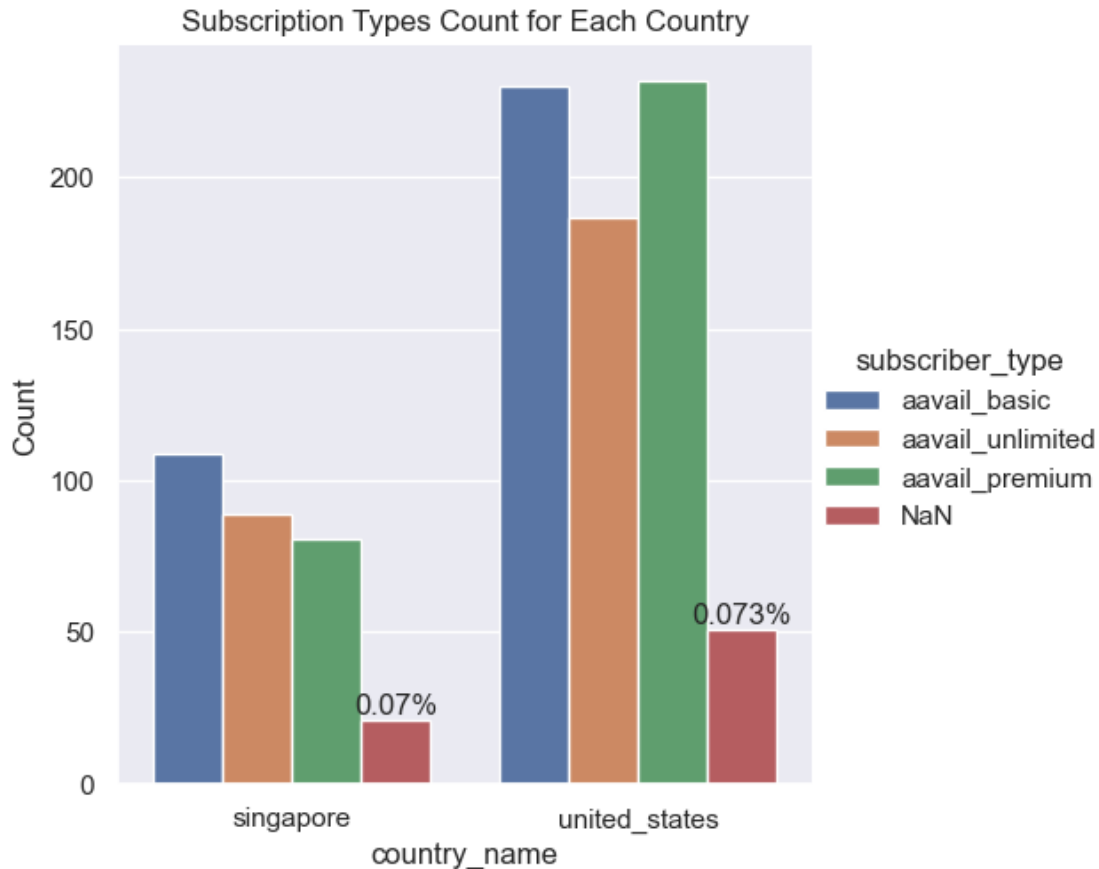
```
[ ]: g = df.groupby(['country_name'])['subscriber_type'].value_counts(dropna=False,
      ↪).to_frame(name='Count').reset_index(level = [0,1]).fillna('NaN')

per1 = g.loc[(g['country_name'] == 'singapore') & (g['subscriber_type'] ==
      ↪'NaN'), 'Count'] / g.loc[g['country_name'] == 'singapore', 'Count'].sum()
per2 = g.loc[(g['country_name'] == 'united_states') & (g['subscriber_type'] ==
      ↪'NaN'), 'Count'] / g.loc[g['country_name'] == 'united_states', 'Count'].sum()

chart = sns.catplot(data = g, kind = 'bar', x = 'country_name', y = 'Count',
      ↪hue='subscriber_type')
chart.set(title='Subscription Types Count for Each Country')
ax = chart.facet_axis(0, 0)
c = ax.containers[3]

# labels = [f'{(v.get_height() / 1000):.1f}K' for v in c]
ax.bar_label(c, labels=[str(round(per1.values[0], 3)) + '%', str(round(per2.
      ↪values[0], 3)) + '%'], label_type='edge')
```

```
[Text(0, 0, '0.07%'), Text(0, 0, '0.073%)]
```



Based on the previous there is no bias based on the **Missing Data**, because the percentages of the missing value are almost the same.

In general there is a bias to **united_states**, so this variable is *unbalanced*.

subscriber_type vs **is_subscriber**

```
[ ]: df.groupby(['subscriber_type'], dropna=False)['is_subscriber'].value_counts().
      to_frame(name='Count').tail(2)
```

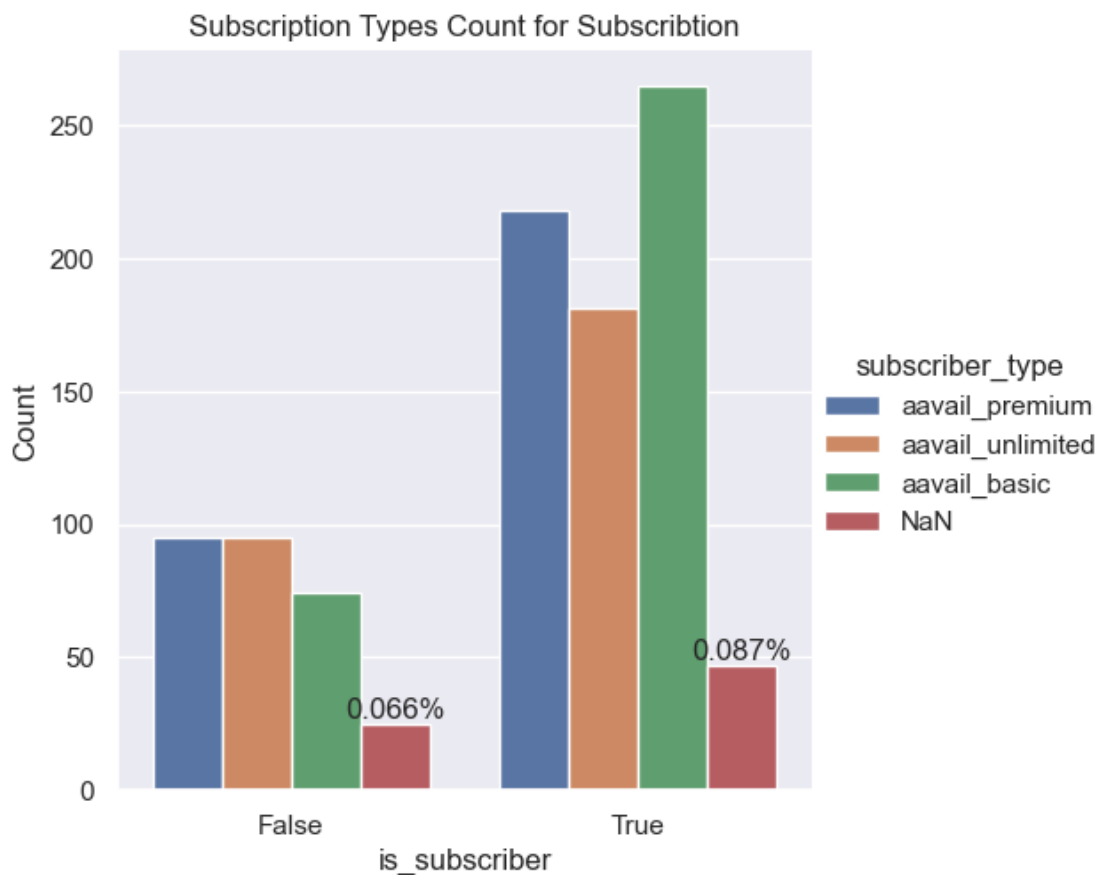
		Count
subscriber_type	is_subscriber	
NaN	True	47
	False	25

```
[ ]: g = df.groupby(['is_subscriber'])['subscriber_type'].value_counts(dropna =
      False).to_frame(name='Count').reset_index(level=[0, 1]).fillna('NaN')

per1 = g.loc[(g['is_subscriber'] == True) & (g['subscriber_type'] == 'NaN'),
      'Count'].values[0] / g.loc[(g['is_subscriber'] == True), 'Count'].sum()
```

```
per2 = g.loc[(g['is_subscriber'] == False) & (g['subscriber_type'] == 'NaN'),  
            ↪ 'Count'].values[0] / g.loc[(g['is_subscriber'] == False), 'Count'].sum()  
  
chart = sns.catplot(data = g, kind = 'bar', x = 'is_subscriber', y = 'Count',  
                    ↪ hue='subscriber_type').set(title='Subscription Types Count for Subscription')  
ax = chart.facet_axis(0, 0)  
c = ax.containers[3]  
  
ax.bar_label(c, labels=[str(round(per1, 3)) + '%', str(round(per2, 3)) + '%'],  
            ↪ label_type='edge')
```

```
[Text(0, 0, '0.066%'), Text(0, 0, '0.087%')]
```



Based on the previous there is no bias based on the **Missing Data**, because the percentages of the missing value are almost the same.

num_streams vs country_name

```
[ ]: df.groupby(['num_streams'], dropna=False)['country_name'].value_counts().  
      ↪ to_frame(name = 'Count').tail(2)
```

		Count
num_streams	country_name	
NaN	united_states	24
	singapore	22

```
[ ]: per1 = df.loc[(df['country_name'] == 'singapore'), 'num_streams'].isna().sum() /
      ↪ len(df.loc[df['country_name'] == 'singapore'])
per2 = df.loc[(df['country_name'] == 'united_states'), 'num_streams'].isna().
      ↪ sum() / len(df.loc[df['country_name'] == 'united_states'])
print(f'singapore missing num of streams: {round(per1, 3)}')
print(f'united_states missing num of streams: {round(per2, 3)}')
```

```
singapore missing num of streams: 0.073
united_states missing num of streams: 0.034
```


The missing data percentage in **singapore** is more than the percentage in **united_states**, so this difference causes bias to the **united_states** data.

We can see that in the next Histogram.

```
[ ]: figure, axes = plt.subplots(1, 2, figsize=(8,5))
figure.suptitle('Histogram for Num of Streams to each Country')
axes[0].set_title('Singapore')
axes[1].set_title('US')


df.loc[(df['country_name'] == 'united_states'), 'num_streams'].
    ↪ plot(kind='hist', bins = 30, ec='black', ax = axes[0])
df.loc[(df['country_name'] == 'singapore'), 'num_streams'].plot(kind='hist',
    ↪ bins = 30, ec='black', color = 'orange', ax = axes[1])

<Axes: title={'center': 'US'}, ylabel='Frequency'>
```



```
[ ]: df['num_streams'].plot(kind='hist', bins = 30, ec='black', title='Number of Streams')
```

```
<Axes: title={'center': 'Number of Streams'}, ylabel='Frequency'>
```

num_streams vs is_subscriber

```
[ ]: df.groupby(['num_streams'], dropna=False)['is_subscriber'].value_counts().
      ↪to_frame(name='Count').tail(2)
```

		Count
num_streams	is_subscriber	
NaN	True	32
	False	14

```
[ ]: per1 = df.loc[(df['is_subscriber'] == True), 'num_streams'].isna().sum() /
      ↪len(df.loc[df['is_subscriber'] == True])
per2 = df.loc[(df['is_subscriber'] == False), 'num_streams'].isna().sum() /
      ↪len(df.loc[df['is_subscriber'] == False])
print(f'Subscribers missing num of streams: {round(per1, 3)}')
print(f'Not Subscribers missing num of streams: {round(per2, 3)}')
```

Subscribers missing num of streams: 0.045

Not Subscribers missing num of streams: 0.048


The Percentage are almost equals so we can say that there is no bias based to the subscribers.

The following is a Histogram for Num of Streams for those who subscribed or not.

```
[ ]: figure, axes = plt.subplots(1, 2, figsize=(8,5))
figure.suptitle('Histogram for Num of Streams to the Subscription')
axes[0].set_title('Yes')
axes[1].set_title('No')

df.loc[(df['is_subscriber'] == True), 'num_streams'].plot(kind='hist', bins = 30, ec='black', ax = axes[0])
df.loc[(df['is_subscriber'] == False), 'num_streams'].plot(kind='hist', bins = 30, ec='black', color = 'orange', ax = axes[1])
```

<Axes: title={'center': 'No'}, ylabel='Frequency'>



We will moving now to Hypotheses Testing, Starting by do *MCAR test* to the **age** and **num_streams**, with the following state of hypothesis:

- \$H_0\$: **age** and **num_streams** are independent so there is no relation between them that effect the missing data.
- \$H_1\$: They are not.

```
[ ]: mt = MCARTest()
value = mt.little_mcar_test( df[['age', 'num_streams']] )
print(f'P-Value: {value}')
```

P-Value: 1.8831016136644507e-05

Based on the P-value we can reject the Null Hypothesis.

state of Hypothesis for **subscriber_type** and **country_name**:

- \$H_0\$: **subscriber_type** and **country_name** are independent so there is no relation between them that effect the missing data.
- \$H_1\$: They are not.

```
[ ]: cc = pd.crosstab(df['subscriber_type'].fillna('missing'), df['country_name'],)
value = stats.chi2_contingency(cc)
print(f'P-Value: {value[1]}')
```

P-Value: 0.26501615774613035

Based on the P-value we can not reject the Null Hypothesis.

state of Hypothesis for **subscriber_type** and **age**: * \$H_0\$: **subscriber_type** and **age** are independent so there is no relation between them that effect the missing data. * \$H_1\$: They are not.

```
[ ]: dfTemp = df[['subscriber_type', 'age']]
dfTemp = dfTemp.fillna('missing')

stats.levene(dfTemp.loc[dfTemp['subscriber_type'] == 'aavail_basic']['age'],
             dfTemp.loc[dfTemp['subscriber_type'] == 'aavail_premium']['age'],
             dfTemp.loc[dfTemp['subscriber_type'] == 'aavail_unlimited']['age'],
             dfTemp.loc[dfTemp['subscriber_type'] == 'missing']['age'],
             center = 'mean')
```

LeveneResult(statistic=2.070637064929318, pvalue=0.10246646882191954)

The groups do not have an equal variance so we can not use *ANOVA*, and instead we will use *Kruskal* which does not have problems with the variance issue.

```
[ ]: dfTemp = df[['subscriber_type', 'age']]
dfTemp = dfTemp.fillna('missing')

stats.kruskal(dfTemp.loc[dfTemp['subscriber_type'] == 'aavail_basic']['age'],
              dfTemp.loc[dfTemp['subscriber_type'] == 'aavail_premium']['age'],
              dfTemp.loc[dfTemp['subscriber_type'] == 'aavail_unlimited']['age'],
              dfTemp.loc[dfTemp['subscriber_type'] == 'missing']['age'],)
```

KruskalResult(statistic=3.9345817178236597, pvalue=0.2686153201366135)

Based on the P-value we can not reject the Null Hypothesis.

```
[ ]: scp.levene(
    df.dropna().loc[df['is_subscriber'] == True]['num_streams'],
    df.dropna().loc[df['is_subscriber'] == False]['num_streams'], center='mean'
)
```

LeveneResult(statistic=20.34474076947048, pvalue=7.341996421949922e-06)

```
[ ]: scp.f_oneway(
    df.dropna().loc[df['is_subscriber'] == True]['num_streams'],
    df.dropna().loc[df['is_subscriber'] == False]['num_streams'],
)
```

F_onewayResult(statistic=19.64331883211847, pvalue=1.0509358853845736e-05)

```
[ ]: scp.levene(
    df.dropna().loc[df['subscriber_type'] == 'aavail_basic']['num_streams'],
    df.dropna().loc[df['subscriber_type'] == 'aavail_premium']['num_streams'],
    df.dropna().loc[df['subscriber_type'] == 'aavail_unlimited']['num_streams'], center='mean'
)
```

LeveneResult(statistic=13.216481438425454, pvalue=2.2114235890901506e-06)

```
[ ]: scp.f_oneway(
    df.dropna().loc[df['subscriber_type'] == 'aavail_basic']['num_streams'],
    df.dropna().loc[df['subscriber_type'] == 'aavail_premium']['num_streams'],
    df.dropna().loc[df['subscriber_type'] == 'aavail_unlimited']['num_streams'],
)
```

F_onewayResult(statistic=1.6117134385494494, pvalue=0.2001338700648013)

```
[ ]: scp.levene(
    df.dropna().loc[df['country_name'] == 'singapore']['num_streams'],
    df.dropna().loc[df['country_name'] == 'united_states']['num_streams'], center='mean'
)
```

LeveneResult(statistic=21.083186117667363, pvalue=5.0372487417334555e-06)

```
[ ]: scp.f_oneway(
    df.dropna().loc[df['country_name'] == 'singapore']['num_streams'],
    df.dropna().loc[df['country_name'] == 'united_states']['num_streams'],
)
```

F_onewayResult(statistic=135.35918442184692, pvalue=3.3231749460711966e-29)

Based on the previous we can say that the missing data type is not **MCAR** and it is almost **MNAR** for **num_streams** variable, and it's **MAR** for **subscriber_type** variable.

Impute **subscriber_type** missing data:

Based on the previous studies we will impute the missing data by sampling from the existense data, because our Missing Data are **MAR** and it's not related to the other features, at the same time it's mode is not make sense to use it because it's almost the same as the others.

```
[ ]: mask = df['subscriber_type'].isna()
temp = df.loc[~mask, ['subscriber_type']].sample(72)['subscriber_type'].values
df.loc[mask, ['subscriber_type']] = temp
df.loc[mask, ['subscriber_type']].head()
```

```
subscriber_type
1  aavail_unlimited
17 aavail_basic
51 aavail_premium
74 aavail_premium
75 aavail_unlimited
```

```
[ ]: df['subscriber_type'].value_counts()
```

```
aavail_basic      372
aavail_premium    331
aavail_unlimited   297
Name: subscriber_type, dtype: int64
```

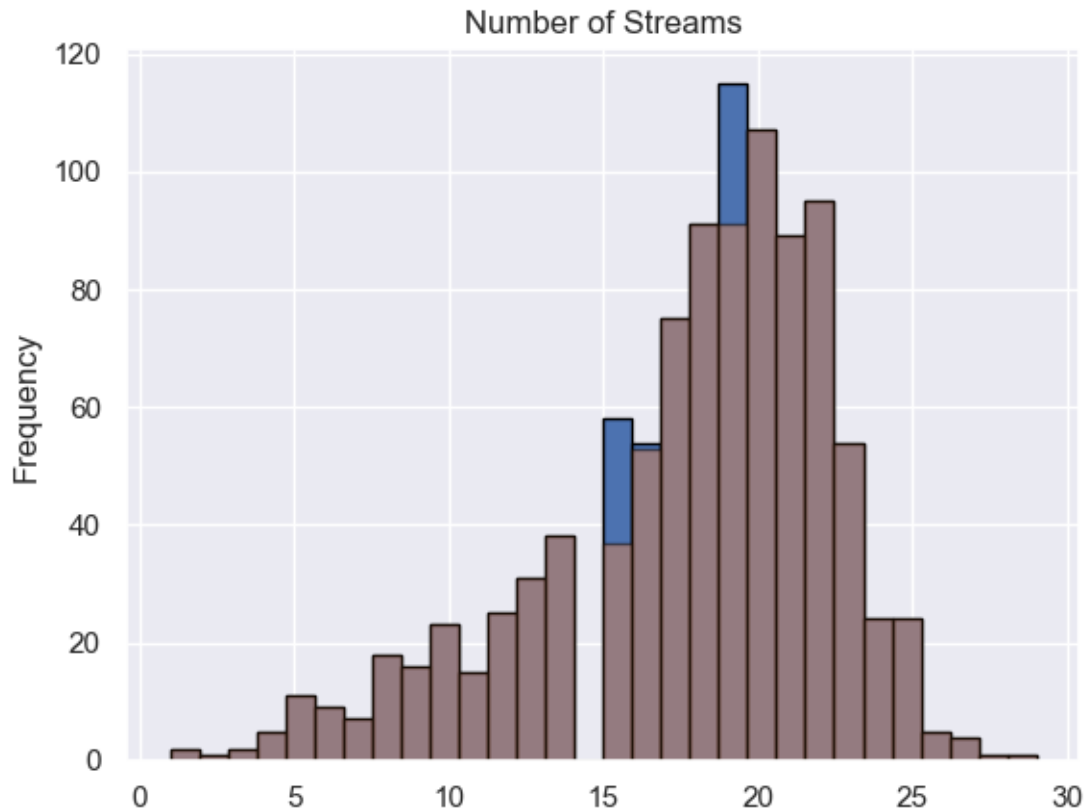
Results: Now we cleared the **subscriber_type** feature from missing data problem.

Impute **num_streams** missing data:

We found in the previous studies that the missing data in **num_streams** feature are **MNAR**, and it's relate to the other features, so we need a model to impute it, and we choose *IterativeImputer* with *BayesianRidge* estimator which is the default option.

```
[ ]: df_impute = pd.concat([df.drop(columns=['country_name']), pd.
    ↳get_dummies(df['country_name'])], axis=1)
imp_num_fet = ['singapore', 'united_states', 'age', 'is_subscriber',
    ↳'num_streams']
imputed = IterativeImputer().fit_transform(df_impute[imp_num_fet])
num_streams = imputed[:,4].round()
```

```
[ ]: df['imputed'] = num_streams
df['imputed'].plot(kind='hist', bins = 30, ec = 'black', title='imputated')
df['num_streams'].plot(kind='hist', bins = 30, ec='black', title='Number of
    ↳Streams', alpha=.5)
df['num_streams'] = num_streams
df.drop(columns=['imputed'], inplace = True)
```




Results: we tried more than one model and the results was the almost the same, now the **num_streams** feature is clear from missing data problem.

1.1.2 EDA

Visualization Box Plots for [**age** | **num_streams**] grouped by **country_name**

```
[ ]: figure, axis = plt.subplots(1, 2, figsize=(16, 5), sharey = True)
box1 = sns.boxplot(df, x = 'age', y = 'country_name', notch=True, orient="h",
    ↪ax= axis[0])
box1.set_title("Box Plot for age Grouped by country_name")

box2 = sns.boxplot(df, x = 'num_streams', y = 'country_name', notch = True,
    ↪orient = 'h', ax = axis[1])
box2.set_title('Box Plot for num_streams Grouped by country_name')
box2.set_ylabel(None)
plt.tight_layout()
```



Results:

- For **age**: we found that there is no difference between the distributions for each country, and there is bias to the left.
- For **num_streams**: we found that there is difference between the distributions and the *united_states*'s is higher.


Bar Plots for [**subscriber_type** | **is_subscriber**] counts grouped by **country_name**

```
[ ]: g = df.groupby(['country_name'])['subscriber_type'].value_counts().
      ↪to_frame(name='Count').reset_index(level=[0,1])
figure, axis = plt.subplots(1, 2, figsize=(16, 5))
bar1 = sns.barplot(g, x = 'country_name', y = 'Count', hue='subscriber_type',
      ↪ax = axis[0])
bar1.set_title('Bar Plot for subscriber_type Counts Grouped by country_name')

g = df.groupby(['country_name'])['is_subscriber'].value_counts().
      ↪to_frame(name='Count').reset_index(level=[0,1])

bar2 = sns.barplot(g, x = 'country_name', y = 'Count', hue = 'is_subscriber',
      ↪ax = axis[1])
bar2.set_title('Bar Plot for is_subscribter Counts Grouped by country_name')
bar2.set_ylabel(None)

plt.tight_layout()
```



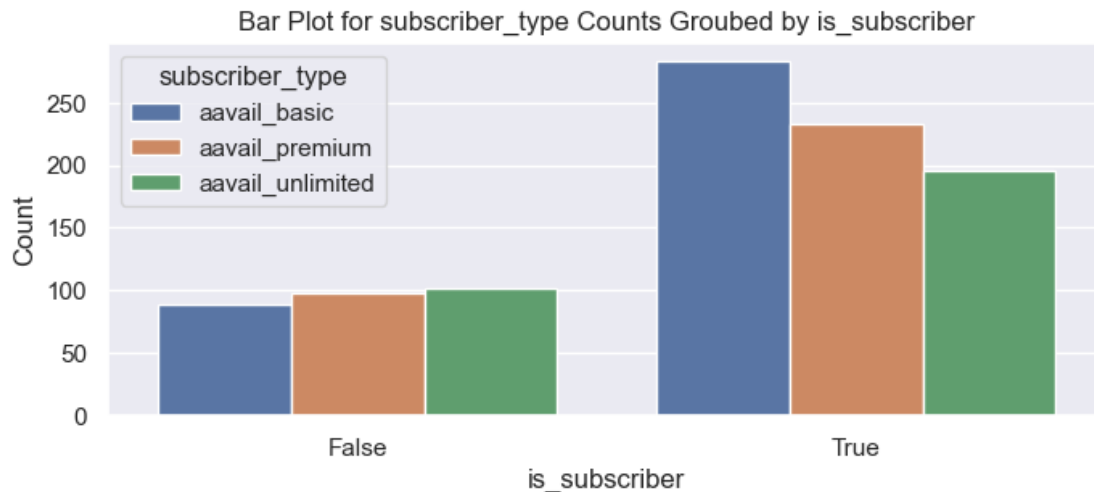
Results:

- For **subscriber_type**: we found that it's almost the same for each country, but *united_states* has more customer than *singapore* by far.
- For **is_subscriber**: we found that in *singapore* the non-subscribers are more than subscribers, but for *united_states* subscribers are the more by far and the number of Non-subscribers are less than the Non-subscribers in *singapore*.

Bar Plot for **subscriber_type** counts grouped by **is_subscriber**

```
[ ]: g = df.groupby(['subscriber_type'])['is_subscriber'].value_counts().  
      ↪to_frame(name='Count').reset_index(level = [0, 1])  
  
plt.figure(figsize = (8, 3))  
bar1 = sns.barplot(g, x = 'is_subscriber', y = 'Count', hue='subscriber_type')  
bar1.set_title('Bar Plot for subscriber_type Counts Groubed by is_subscriber')
```

Text(0.5, 1.0, 'Bar Plot for subscriber_type Counts Groubed by is_subscriber')



Results:

- Based on the **subscriber_type** the Non-subscribers count are almost the same.

Box Plots for [**age** | **num_streams**] grouped by subscriber_type


```
[ ]: figure, axis = plt.subplots(1, 2, figsize=(16, 5), sharey = True)  
box1 = sns.boxplot(df, x = 'age', y = 'subscriber_type', notch=True,   
      ↪orient="h", ax= axis[0])  
box1.set_title("Box Plots for age grouped by subscriber_type")
```



```

box2 = sns.boxplot(df, x = 'num_streams', y = 'subscriber_type', notch = True,
    ↪orient = 'h', ax = axis[1])
box2.set_title('Box Plots for num_streams grouped by subscriber_type')
box2.set_ylabel(None)
plt.tight_layout()

```



Results:

- For **age**: the distributions has almost the same midians, just the unlimited type has bigger distribution, and it is skewed to the left.
- For **num_streams**: they all similar, with a little bit difference with the median in the unlimited type.

Box Plots for [**age** | **num_streams**] grouped by **country_name**, grouped by ***subscriber_type**


```

[ ]: figure, axis = plt.subplots(1, 2, figsize = (16, 5), sharey = True)
box1 = sns.boxplot(df, y = 'country_name', x = 'age', hue='subscriber_type',
    ↪ax=axis[0])
box1.set_title('Box Plots for age grouped by country_name, grouped by
    ↪subscriber_type')

box2 = sns.boxplot(df, y = 'country_name', x = 'num_streams',
    ↪hue='subscriber_type', ax=axis[1])
box2.set_title('Box Plots for num_streams grouped by country_name, grouped by
    ↪subscriber_type')
box2.set_ylabel(None)

plt.tight_layout()

```



Results:

- For **age** and *united_states*: the unlimited type has bigger distribution, but they all have almost the same median.
- For **age** and *singapore*: the premium type has the bigger distribution then the basic and the last one is the unlimited, and they all have similar medians.
- For **num_streams** and *united_states*: they almsot similar, the premium type has bigger median.
- For **num_streams** and *singapore*: there is difference between types, the basic type has smaller distribution but higher values and higher median.

Scatter Plots for [**age** | **num_streams**] vs **subscriber_type** separated by **is_subscriber**

```
[ ]: figure, axis = plt.subplots(1, 2, figsize = (16, 6), sharey = True)

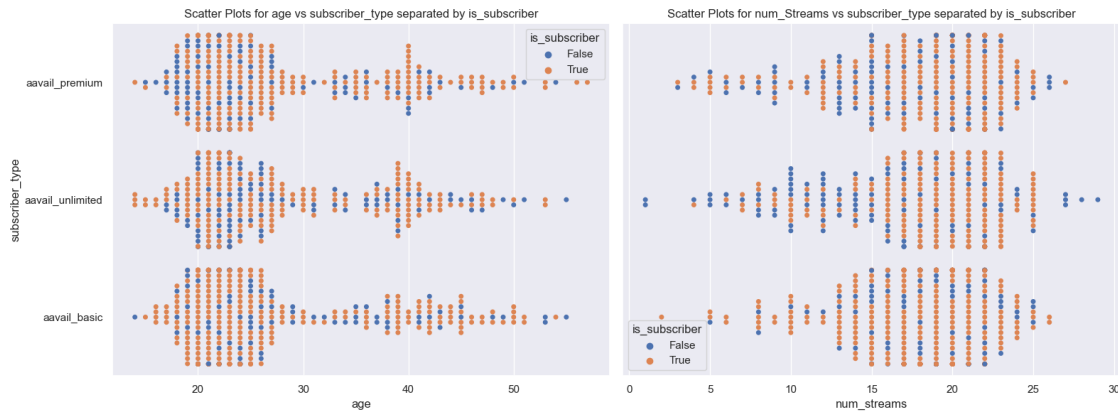
sca1 = sns.swarmplot( data=df, x = 'age', y = 'subscriber_type', hue = 'is_subscriber', ax = axis[0])
sca1.set_title('Scatter Plots for age vs subscriber_type separated by is_subscriber')

sca2 = sns.swarmplot(data = df, x = 'num_streams', y = 'subscriber_type', hue = 'is_subscriber', ax = axis[1])
sca2.set_title('Scatter Plots for num_Streams vs subscriber_type separated by is_subscriber')
sca2.set_ylabel(None)

plt.tight_layout()
```

```
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 6.9% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 16.4% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
warnings.warn(msg, UserWarning)
```

```
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 16.9% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 11.1% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 27.4% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 28.4% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 18.2% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 38.7% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 13.0% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 8.4% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 23.1% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 24.2% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 13.5% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
c:\Users\ASUS\anaconda3\lib\site-packages\seaborn\categorical.py:3544:
UserWarning: 34.4% of the points cannot be placed; you may want to decrease the
size of the markers or use stripplot.
    warnings.warn(msg, UserWarning)
```




Results:

- For **age**: we found that the Non-subscribers with premium type are centered around the age of 20, and around 23 with the unlimited type.
- For **num_streams**: we found that people with higher and lower number of streams have the highest number of Non-subscribers.

Scatter and density Plots for **age** vs **num_streams** separated by **is_subscriber**

```
[ ]: sns.pairplot(df[['age', 'num_streams', 'is_subscriber']], hue='is_subscriber',
↪corner=True, height=3)
```

<seaborn.axisgrid.PairGrid at 0x21a31e70190>



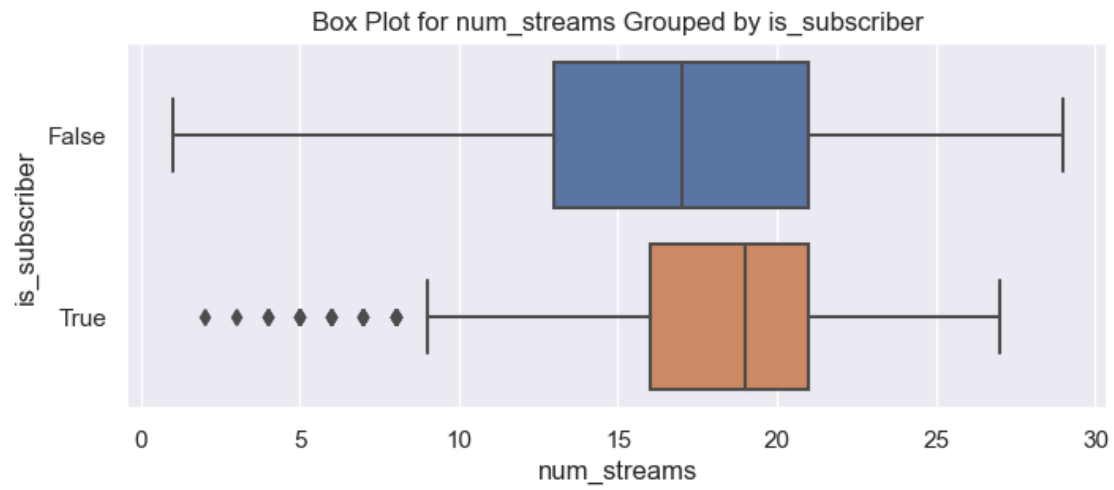
Results:

- There is no clear relation between **age** and **num_streams**.

Box Plot for **num_streams** grouped by **is_subscriber**

```
[ ]: plt.figure(figsize=(8, 3))
      box = sns.boxplot(data=df, x = 'num_streams', y = 'is_subscriber', orient='h')
      box.set_title('Box Plot for num_streams Grouped by is_subscriber')
```

Text(0.5, 1.0, 'Box Plot for num_streams Grouped by is_subscriber')



Results:

- Non_subscribers have bigger distribution than subscribers, but the median is lower.